Lethal Autonomous Weapon Systems: Ethical Perspectives

Group 2: Erika Gallegos, Dylan Scott-Dawkins, Rogelio Aguilar

AAI 531-003 | Spring 2026 | Final Project

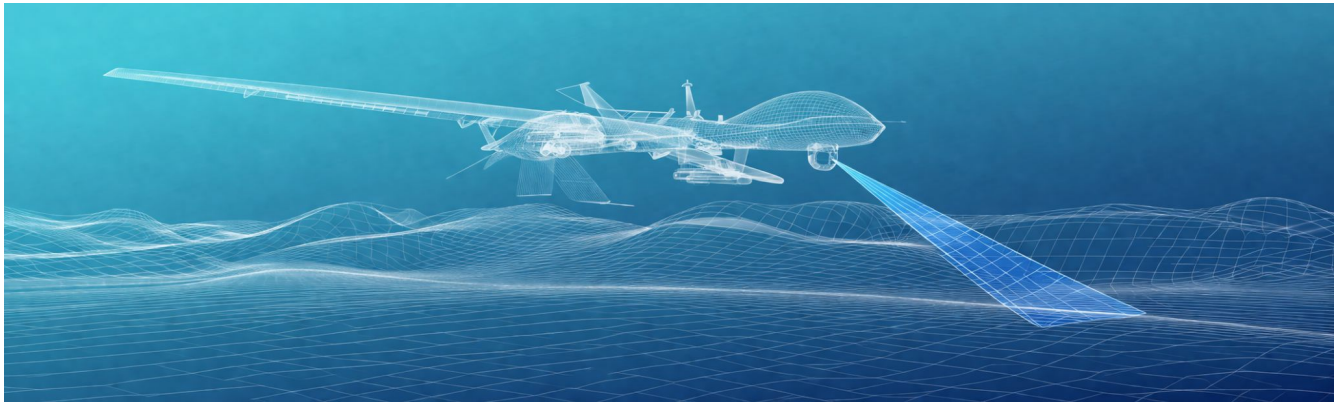

AI Ethics Principles

- **Fairness:** Lethal autonomous weapons systems may struggle to distinguish between combatants and civilians, increasing the risk of biased or unjust outcomes due to flawed training data or limited contextual understanding (Russell et al., 2015).

- **Accountability:** The use of autonomous and swarm based systems creates an “accountability gap,” making it difficult to determine whether engineers, operators, or policymakers are liable when harm occurs (Scharre, 2018).

- **Transparency:** The opaque nature of AI decision making, especially in swarm intelligence, limits explainability, undermining trust and oversight. Governments should provide partial transparency (e.g., safety standards, oversight mechanisms) while protecting sensitive capabilities.

- **Key Concern:** The combination of low explainability and rapid decision making challenges human control, ethical judgment, and compliance with international law (Viveros Alvarez, 2022).

Unintended Stakeholders

- **Civilian populations in conflict zones**: most directly at risk from misidentification
- **Military personnel and commanders:** face legal and moral liability for actions taken by systems they may not fully understand or be able to override in time
- **International humanitarian organizations**: must adapt legal frameworks (e.g., Geneva Conventions) to address weapons that operate without human intent behind each use of force

Unintended Impacts

- **Escalation risk**: AI systems in simulated wargames have frequently chosen nuclear escalation, suggesting autonomous decision-making may amplify rather than contain conflict
- **Accountability vacuum**: No clear legal framework exists to assign blame when an autonomous system causes civilian casualties - the "responsibility gap" (Scharre, 2018)
- **Automation bias**: Human operators may over-trust AI recommendations, rubber-stamping targeting decisions rather than exercising independent judgment

Erika’s Perspective

- Removes humans from the **moral consequences** of action
  - Increases likelihood of escalation
- Proponents argue these systems protect warfighters by distancing them from conflict
  - However, **human-in-the-loop** drones already achieve this
  - Full autonomy is not necessary for safety
- Delegating lethal decisions to machines creates an **accountability gap**
- Ultimately, computers should not decide who lives or dies

Dylan’s Perspective

- I see this as fundamentally a **trust and transparency problem**. These systems are built on machine learning models whose decision logic is opaque even to their designers. Asking a commander to take legal and moral responsibility for a targeting decision made by a system they cannot explain contradicts basic principles of accountability
- **My common-sense view**: I believe that decisions about life and death require moral reasoning, contextual judgment, and the weight of personal responsibility: qualities that machines do not possess. Speed and efficiency should never outweigh the ethical obligation to ensure a human being genuinely decides when lethal force is used.

Rogelio’s Perspective

- While autonomous weapons may offer efficiency, I want to focus on the ethical risks. These systems challenge fairness by potentially misidentifying targets, lack accountability when things go wrong, and operate with limited transparency. As a result, without human judgment, they cannot fully interpret context, making their use in warfare deeply concerning.
- My perspective is grounded in a deontological framework, which emphasizes duties and moral rules rather than outcomes. Autonomous weapons are problematic because they risk violating core duties like protecting human dignity and ensuring accountability. If an AI system makes an irreversible decision without human judgment or clear responsibility, it becomes difficult to justify morally. Additionally, the lack of transparency means these decisions cannot be fully explained, which further conflicts with ethical standards. Therefore, I believe lethal autonomous weapon systems should not be used in warfare.

Conclusion

- Lethal autonomous weapons create ethical risks around fairness, accountability, and transparency
- Responsibility gap undermines clear moral and legal ownership of harm
- AI lacks the human judgement required for life-and-death decisions
- **Meaningful human control must remain central to any use of lethal force**